


Input & Output


1. I/O Management I

1.1 CSCI 330

CSCI 330 UNIX and Network Programming



I/O Management I



1.2 Unit Overview

Unit Overview

- System calls
- File I/O
 - open, creat
 - read, write
 - close
- more:
 - unlink, stat, chmod, dup

1.3 UNIX System Call

UNIX System Call

- a system call is how a program requests a service from the operating system, i.e. UNIX kernel
- system call executes code in the kernel and makes direct use of facilities provided by the kernel

versus:

- library function is linked to executable, becomes part of the executable

1.4 System Call Categories

System Call Categories

- file management
 - create/delete file, open/close, read/write, get/set attributes
- process control
 - create/terminate process, wait/signal event, allocate/free memory
- communication
 - create/delete connection, send/receive messages, remote devices
- information management
 - get/set system data and attributes, e.g. time
- device management
 - attach/request/release/detach device, read/write/position

1.5 System Call Invocation

System Call Invocation

- declare system call via appropriate C header file
 - read man page carefully (section 2 of manual)
- prepare parameters using basic C data types
 - no C++ classes, i.e. string
- prepare suitable return value variable

call like any other function

1.6 File Management

File Management

open	open a file
read	read data from a file
write	write data to a file
close	close a file
creat	make a new file

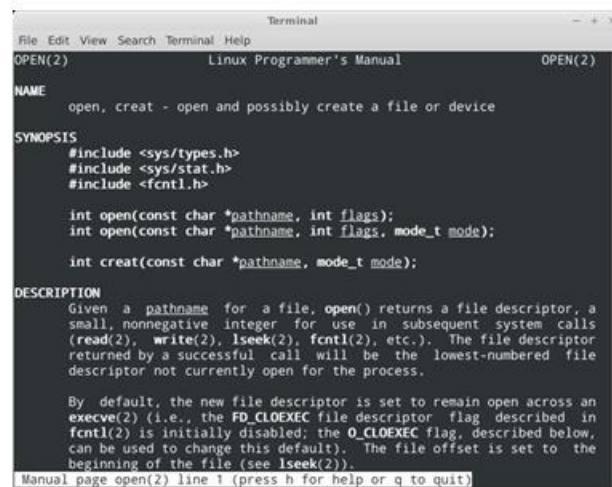
- next:

unlink	remove file
stat	get file information
chmod	change permissions
dup	duplicate file descriptor

- all calls share file descriptor, i.e. number, to identify file

1.7 System Call: open

System Call: open



```
Terminal
File Edit View Search Terminal Help
OPEN(2) Linux Programmer's Manual OPEN(2)

NAME
open, creat - open and possibly create a file or device

SYNOPSIS
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int open(const char *pathname, int flags);
int open(const char *pathname, int flags, mode_t mode);
int creat(const char *pathname, mode_t mode);

DESCRIPTION
Given a pathname for a file, open() returns a file descriptor, a
small, nonnegative integer for use in subsequent system calls
(read(2), write(2), lseek(2), fcntl(2), etc.). The file descriptor
returned by a successful call will be the lowest-numbered file
descriptor not currently open for the process.

By default, the new file descriptor is set to remain open across an
execve(2) (i.e., the FD_CLOEXEC file descriptor flag described in
fcntl(2) is initially disabled; the O_CLOEXEC flag, described below,
can be used to change this default). The file offset is set to the
beginning of the file (see lseek(2)).

Manual page open(2) line 1 (press h for help or q to quit)
```

1.8 System Call: open

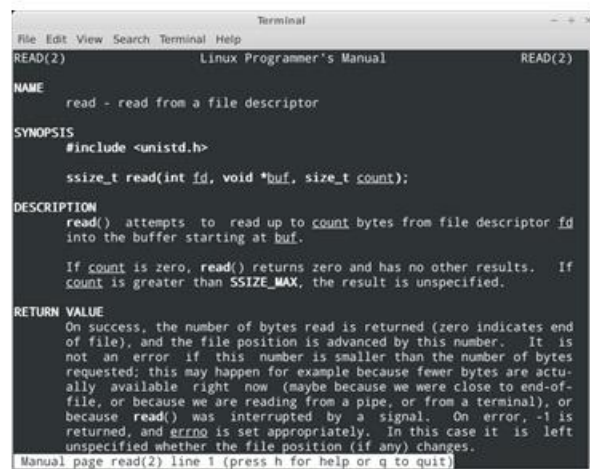
System Call: open

```
int open(const char* pathname, int flags)
```

- opens file specified as **pathname** for access
- **flags** determine access type
 - `O_RDONLY` read only
 - `O_WRONLY` write only
 - `O_RDWR` read and write
- returns file descriptor, to be used in read/write/close
- returns -1 on error

1.9 System Call: read

System Call: read



1.10 System Call: read

System Call: read

```
ssize_t read(int fd, void *buf, size_t count)
```

- attempts to read **count** bytes from file descriptor **fd** into the buffer starting at **buf**
 - `ssize_t` and `size_t` are system specific integer types
- returns the number of bytes read
 - maybe smaller than count, zero indicates end of file
 - file position is advanced by this number
- returns -1 on error

1.11 System Call: close

System Call: close

```
int close(int fd)
```

- closes file specified by **fd** file descriptor
 - makes file descriptor available
- returns zero on success

1.12 System Call: read example

System Call: read example

```
read.cxx - /home/student/Desktop/Unit 12 - IO Mgmt - Geany
File Edit Search View Document Project Build Tools Help

read.cxx
18 int main() {
19     int fd, count;
20     char buffer[1024], filename[] = "example.txt";
21     // open file
22     fd = open(filename, O_RDONLY);
23     if (fd == -1) {
24         perror(filename);
25         exit(EXIT_FAILURE);
26     }
27     // read from file
28     count = read(fd, buffer, 1024);
29     if (count == -1) {
30         perror(filename);
31         exit(EXIT_FAILURE);
32     }
33     cout << "read " << count << " bytes from file\n";
34     // close file
35     close(fd);
36     return 0;
37 }
```

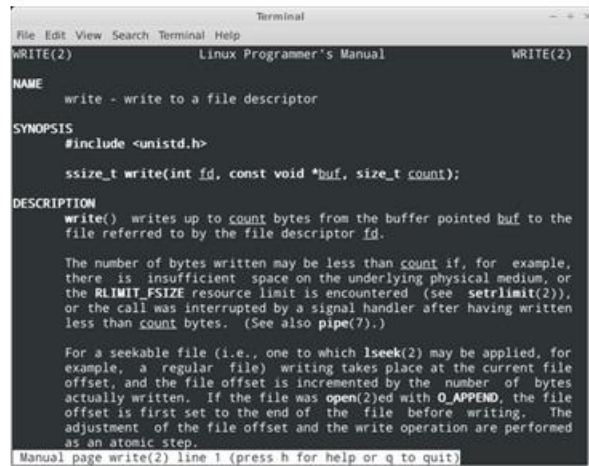
1.13 System Call: readAll example

System Call: readAll example

```
readAll.cxx
18 int main(int argc, char* argv[]) {
19     int fd, count, sum=0;
20     // check command line arguments
21     if (argc < 2) {
22         cerr << "Usage: readAll filename\n";
23         exit(EXIT_FAILURE);
24     }
25     char buffer[1024];
26     // open file
27     fd = open(argv[1], O_RDONLY);
28     if (fd == -1) {
29         perror(argv[1]);
30         exit(EXIT_FAILURE);
31     }
32     // read from file
33     while ((count = read(fd, buffer, 1024)) != 0) {
34         if (count == -1) {
35             perror(argv[1]);
36             exit(EXIT_FAILURE);
37         }
38         sum += count;
39         cout << "read " << count << " bytes from file\n";
40     }
41     cout << "read all " << sum << " bytes from " << argv[1] << endl;
42     // close file
43     close(fd);
44     return 0;
45 }
```

1.14 System Call: write

System Call: write



```
Terminal
File Edit View Search Terminal Help
WRITE(2) Linux Programmer's Manual WRITE(2)

NAME
    write - write to a file descriptor

SYNOPSIS
    #include <unistd.h>

    ssize_t write(int fd, const void *buf, size_t count);

DESCRIPTION
    write() writes up to count bytes from the buffer pointed buf to the
    file referred to by the file descriptor fd.

    The number of bytes written may be less than count if, for example,
    there is insufficient space on the underlying physical medium, or
    the RLIMIT_FSIZE resource limit is encountered (see setrlimit(2)),
    or the call was interrupted by a signal handler after having written
    less than count bytes. (See also pipe(7).)

    For a seekable file (i.e., one to which lseek(2) may be applied, for
    example, a regular file) writing takes place at the current file
    offset, and the file offset is incremented by the number of bytes
    actually written. If the file was open(2)ed with O_APPEND, the file
    offset is first set to the end of the file before writing. The
    adjustment of the file offset and the write operation are performed
    as an atomic step.

Manual page write(2) line 1 (press h for help or q to quit)
```

1.15 System Call: write

System Call: write

`ssize_t write(int fd, const void *buf, size_t count)`

- writes up to **count** bytes from buffer starting at **buf** to the file referred to by file descriptor **fd**
 - `ssize_t` and `size_t` are system specific integer types
- returns the number of bytes written
 - maybe smaller than count
- returns -1 on error

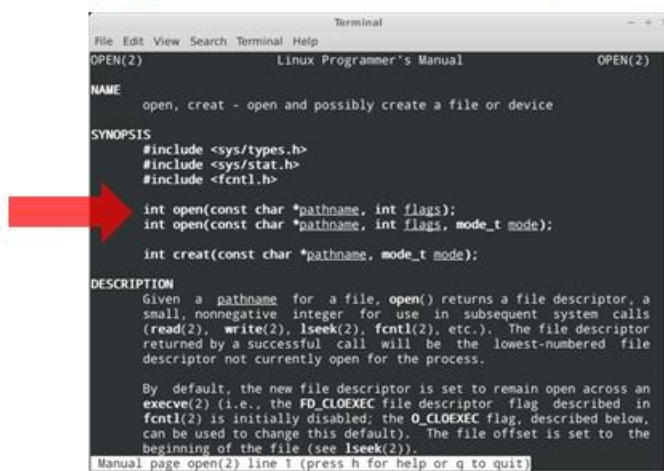
1.16 System Call: write example

System Call: write example

```
write.cxx
17 int main() {
18     int fd, count;
19     char buffer[] = "one line to write";
20     // open existing file, will overwrite current content
21     fd = open("example.txt", O_WRONLY);
22     if (fd == -1) {
23         perror("example.txt");
24         exit(fd);
25     }
26     // write to file
27     count = write(fd, buffer, sizeof(buffer));
28     if (fd == -1) {
29         perror("example.txt");
30         exit(fd);
31     }
32     cout << "wrote " << count << " bytes to file\n";
33     // close file
34     close(fd);
35     return 0;
36 }
```

1.17 revisit System Call: open

revisit System Call: open



```
Terminal
File Edit View Search Terminal Help
OPEN(2) Linux Programmer's Manual OPEN(2)

NAME
open, creat - open and possibly create a file or device

SYNOPSIS
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int open(const char *pathname, int flags);
int open(const char *pathname, int flags, mode_t mode);
int creat(const char *pathname, mode_t mode);

DESCRIPTION
Given a pathname for a file, open() returns a file descriptor, a
small, nonnegative integer for use in subsequent system calls
(read(2), write(2), lseek(2), fcntl(2), etc.). The file descriptor
returned by a successful call will be the lowest-numbered file
descriptor not currently open for the process.

By default, the new file descriptor is set to remain open across an
execve(2) (i.e., the FD_CLOEXEC file descriptor flag described in
fcntl(2) is initially disabled; the O_CLOEXEC flag, described below,
can be used to change this default). The file offset is set to the
beginning of the file (see lseek(2)).

Manual page open(2) line 1 (press h for help or q to quit)
```

1.18 System Call: open

System Call: open

```
int open(const char* pathname, int flags)
```

- opens file specified as **pathname** for access
- **flags** determine access type
 - **O_RDONLY** read only
 - **O_WRONLY** write only
 - **O_RDWR** read and write
- returns file descriptor, to be used in read/write/close
- returns -1 on error

1.19 System Call: open

System Call: open

- additional **flags**, used with **O_WRONLY**:
 - **O_APPEND** to append to an existing file
 - **O_TRUNC** existing file will overwritten (default)

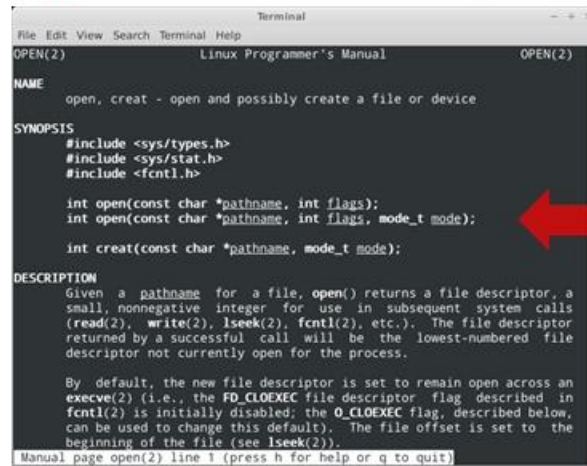
Ex.:

```
O_WRONLY | O_TRUNC  
O_WRONLY | O_APPEND
```

- additional flag: **O_CREAT**
 - creates file, if file does not exist

1.20 revisit System Call: open

revisit System Call: open



```
Terminal
File Edit View Search Terminal Help
Linux Programmer's Manual
OPEN(2)

NAME
    open, creat - open and possibly create a file or device

SYNOPSIS
    #include <sys/types.h>
    #include <sys/stat.h>
    #include <fcntl.h>

    int open(const char *pathname, int flags);
    int open(const char *pathname, int flags, mode_t mode);
    int creat(const char *pathname, mode_t mode);

DESCRIPTION
    Given a pathname for a file, open() returns a file descriptor, a
    small, nonnegative integer for use in subsequent system calls
    (read(2), write(2), lseek(2), fcntl(2), etc.). The file descriptor
    returned by a successful call will be the lowest-numbered file
    descriptor not currently open for the process.

    By default, the new file descriptor is set to remain open across an
    execve(2) (i.e., the FD_CLOEXEC file descriptor flag described in
    fcntl(2) is initially disabled; the O_CLOEXEC flag, described below,
    can be used to change this default). The file offset is set to the
    beginning of the file (see lseek(2)).

Manual page open(2) line 1 (press h for help or q to quit)
```

1.21 System Call: open with O_CREAT

System Call: open with O_CREAT

```
int open(const char* pathname, int flags, mode_t mode)
```

- must specify file mode, i.e. permissions, of type `mode_t`

```
S_IRWXU  00700  user has read, write and execute permission
S_IRUSR  00400  user has read permission
S_IWUSR  00200  user has write permission
S_IXUSR  00100  user has execute permission
S_IRWXG  00070  group has read, write and execute permission
S_IRWXO  00007  others have read, write and execute perm.
...
```

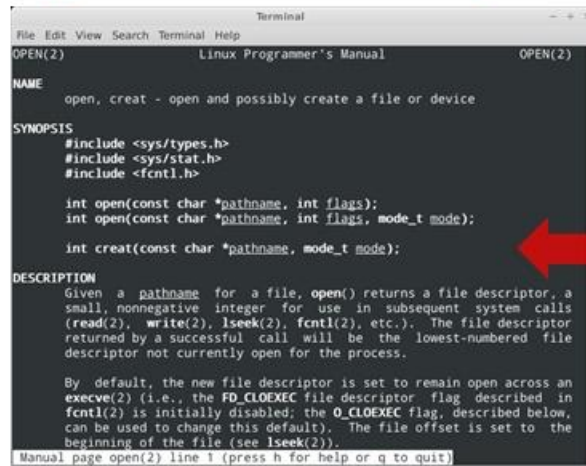
- C++ requires leading '0' for octal numbers
- actual file mode is further modified by umask

Example:

```
open("ex.txt", O_WRONLY | O_APPEND | O_CREAT, 00666)
```

1.22 revisit System Call: open

revisit System Call: open



```
File Edit View Search Terminal Help
OPEN(2)      Linux Programmer's Manual      OPEN(2)

NAME
  open, creat - open and possibly create a file or device

SYNOPSIS
  #include <sys/types.h>
  #include <sys/stat.h>
  #include <fcntl.h>

  int open(const char *pathname, int flags);
  int open(const char *pathname, int flags, mode_t mode);
  int creat(const char *pathname, mode_t mode);

DESCRIPTION
  Given a pathname for a file, open() returns a file descriptor, a
  small, nonnegative integer for use in subsequent system calls
  (read(2), write(2), lseek(2), fcntl(2), etc.). The file descriptor
  returned by a successful call will be the lowest-numbered file
  descriptor not currently open for the process.

  By default, the new file descriptor is set to remain open across an
  execve(2) (i.e., the FD_CLOEXEC file descriptor flag described in
  fcntl(2) is initially disabled; the O_CLOEXEC flag, described below,
  can be used to change this default). The file offset is set to the
  beginning of the file (see lseek(2)).

Manual page open(2) line 1 (press h for help or q to quit)
```

1.23 System Call: creat

System Call: creat

```
int creat(const char *pathname, mode_t mode)
```

- creates new file specified as **pathname** and opens file for write access
- **mode** specifies permissions of type **mode_t**
- returns file descriptor
- returns -1 on error

1.24 Example: copy file

Example: copy file

```
copy.cxx
1  /*
2   * copy.cxx
3   *
4   * Example Program for CSCI 330
5   * implement copy by reading file, and writing file
6   * via read/write system calls
7   *
8   */
9  #include <sys/types.h>
10 #include <sys/stat.h>
11 #include <fcntl.h>
12 #include <unistd.h>
13 #include <stdlib.h>
14 #include <stdio.h>
15 #include <iostream>
16 using namespace std;
17
18 int main(int argc, char* argv[]) {
19     int ifd, ofd, count, sum=0;
20     char buffer[1024];
21
22     // check command line arguments
23     if (argc < 3) {
24         cerr << "Usage: copy fromFile toFile\n";
25         exit(EXIT_FAILURE);
26     }
27     // open file to read
28     ifd = open(argv[1], O_RDONLY);
29     if (ifd == -1) {
30         perror(argv[1]);
31         exit(EXIT_FAILURE);
32     }
33 }
```

1.25 Summary

Summary

- System calls: open, creat, read, write, close
- next:
 - unlink remove file
 - stat get file information
 - chmod change permissions
 - dup duplicate file descriptor